

TCS^AI GOOGLE HACKATHON 2026 · TRACK 1: INTELLIGENT SEARCH & GEMINI
ENTERPRISE



Regulatory Radar **Impact**

An automated regulatory nervous system that scans for regulation changes, parsing them with Semantic Knowledge Graphs to instantly map compliance gap impacts across the enterprise.

ARCHITECTURE

Multi-Agent Routing

MODEL

Gemini Enterprise Plus (Gemini 1.5 Pro)

KNOWLEDGE BASE

Google Drive Connector & Files

Sana Iqbal

Employee ID: 1213383

SECTION 01

Executive Summary

20% Weight

 **The Organizational Problem:** Assessing new regulatory changes requires compliance teams to manually read hundreds of pages to locate which specific internal IT policies, SLA thresholds, and Cloud architectures require change. The sheer manual effort causes delayed compliance reporting and introduces human error on penalty-driving regulations.

 **Our Solution: Regulatory Radar** acts as a proactive multi-agent auditor. Using Vertex AI to map regulatory amendments to an internal Google Drive Knowledge connector indexing past internal documentation, it actively maps out internal policies requiring updates and flags potential gaps autonomously in seconds.

What's Unique: Regulatory Radar generates **structured Impact Cards** (severity, gap, evidence citations, owners, deadlines) by grounding regulatory clause deltas against a **Google Drive Knowledge Connector**. The multi-agent design separates "clause extraction" from "impact mapping," reducing hallucinations and producing repeatable, auditable outputs.

SECTION 02

Multi-Agent Architecture

15% Weight

1. Ideathon Concept Blueprint

Regulatory Change Impact Radar - Architecture Diagram




Figure 1a: Conceptual Design — Pipeline connecting external regulatory sources, Document AI parsers, and Vertex AI Semantic Knowledge Graphs

2. Vertex AI Live Implementation



Figure 1b: Live Implementation — The Orchestrator routes parsing tasks to the Regulatory Analyzer, then maps the results to internal policies using the Impact Mapper




Figure 1c: Regulatory Radar native Agent UI interface displaying conversation starters upon user entry




Figure 1d: Agent Designer UI displaying the interconnected Multi-Agent architecture




Figure 1e: Core agent model selection mapping to Gemini Enterprise Plus

SECTION 03

Core Prompt Engineering & Guardrails

15% Weight

Stringent prompt formatting and robust instruction layers allow the Regulatory Radar agents to output uniform evaluation metric cards containing explicit document path citations (Grounding) and exact severity warnings.




Figure 2a: Multi-agent definitions outlining clear role boundaries and rigid UI constraint enforcement




Figure 2b: Zero-Hallucination rules, rigorous PII prevention blocks, and Human-In-The-Loop disclaimers

SECTION 04

Enterprise Grounding

20% Weight

The system evaluates incoming regulatory mandates strictly against 16 core internal documents natively integrated into a knowledge connector via Google Drive. Grounding ensures high accuracy mapping against CRM systems, incident runbooks, and existing data retention policies.




Figure 3a: Enterprise Google Drive integrated in Vertex AI Knowledge Config via Connector




Figure 3b: Source directory holding internal IT policies, vendor agreements, and breach timelines

SECTION 05

Proof of Execution: DPDPA Impact

20% Weight

By processing a hypothetical new amendment to the DPDPA, we validate the agent's capability to read text inputs and seamlessly synthesize them against internal documentation policies.




Figure 4a: Agent outputs High Risk notifications with rigid evaluation parameters




Figure 4b: Explicit Gap mappings identifying required changes dynamically referencing internal policies




Figure 4c: Overall gap and system breakdown with explicitly grounded citations and a generated Action Plan

SECTION 05 (CONTINUED)

Compliance Evaluation & Scheduling

The system is able to evaluate real-world frameworks and deliver robust compliance scorecards summarizing organizational readiness parameters and overall audit capabilities.




Figure 4d: Summary evaluation on SEBI Framework showcasing exact areas passing and failing audit validations




Figure 4e: Continued documentation of individual criteria compliance mapping across vendor agreements




Figure 4f: Regulatory Landscape Onboarding summaries providing rapid ramp up to new compliance officers covering core business policies

SECTION 06

Continuous Execution Scheduling

10% Weight




Figure 5a: Regulatory sweep recurring background schedule parameter definitions to automatically invoke sweeps

SECTION 07

Scalability, Business Value & Responsible AI 10% Weight

Enterprise Scalability Roadmap

Phase	Timeline	Scope	Google Cloud Services
PoC (Current)	2–3 weeks	3 agents, 16 internal docs, single regulatory domain	Vertex AI Agent Builder, Gemini Enterprise Plus, Google Drive Connector
MVP	4–6 weeks	10+ regulatory frameworks, multi-department coverage, scheduled sweeps	+ Cloud Scheduler, BigQuery audit logging, Pub/Sub notifications
Production	8–12 weeks	100+ regulations, real-time legislative monitoring, cross-geography compliance	+ Cloud Functions, Document AI, Vertex AI Search, multi-region deployment

Quantitative Business Impact

90% FASTER IMPACT ANALYSIS	16 POLICY DOCS ANALYZED	3 SPECIALIZED AGENTS	100% GROUNDED CITATIONS
--------------------------------------	-----------------------------------	--------------------------------	-----------------------------------

💰 **Strategic Value for TCS:** Regulatory Radar transforms TCS into a **proactive compliance partner** for BFSI, Healthcare, and Insurance clients. Instead of reactive audit scrambles, clients get instant gap analysis and action plans whenever a regulation changes — reducing compliance cycle time from **weeks to minutes** and eliminating penalty-driving human error.

Security & Responsible AI

Aspect	Implementation
Data Privacy	System instructions explicitly block PII extraction. Agent refuses to process personal data in regulatory queries.
Hallucination Prevention	All outputs grounded via Google Drive Knowledge Connector — agent cites specific source documents.
Human-in-the-Loop	Every assessment includes mandatory disclaimers: recommendations require human review before implementation.
Access Control	Agent access scoped to read-only connector. No write permissions to enterprise documentation.
Synthetic Data	All PoC regulatory documents and policies are synthetic — no real client or regulatory data.

SECTION 08

Conclusion & Evidence Checklist

Regulatory Change Impact Radar completely mitigates blind spots in enterprise compliance updating architectures by linking directly against internal documentation using a high-density, strictly formatted multi-agent evaluation hierarchy containing fully compliant PII and Audit guardrails.

✔ Evidence Checklist

✔ 1c. Welcome Interface Displayed	✔ 1d. Orchestrator Flow Canvas View
✔ 1e. Agent Engine Configuration	✔ 2a. System Instructions Structure
✔ 2b. Mandatory PII/HITL Guardrails	✔ 3a. Knowledge Connector Integration
✔ 3b. Enterprise Google Drive Memory	✔ 4a. DPDPA Severity Assessment
✔ 4b. Gap Mappings Validated	✔ 4c. Action Plan & Citations
✔ 4d. SEBI Framework Scorecard	✔ 4f. Officer Onboarding Summaries
✔ 5a. Automatic Sweeps Configured	

APPENDIX A

Agent System Instructions & Security Guardrails

ORCHESTRATOR AGENT — SYSTEM INSTRUCTIONS

Routing Logic & Delegation Rules

You are the "Regulatory Radar Orchestrator" – the central routing agent for the Regulatory Change Impact Radar multi-agent system.

Your **ONLY** job is to determine which specialist agent should handle the user's request and delegate the task to them. You **MUST NOT** answer the user's compliance or regulatory questions yourself.

YOUR SPECIALIST AGENTS

AGENT 1 – "Regulatory Analyzer"

- TRIGGER: User submits a new regulatory document, circular, or legal text for analysis.
- TRIGGER: User asks to summarize or extract requirements from a regulation.

AGENT 2 – "Impact Mapper"

- TRIGGER: User asks about the **IMPACT** of a regulation on the enterprise.
- TRIGGER: User asks "Are we compliant with X?" or "What needs to change?"

BEHAVIORAL RULES & GUARDRAILS (MANDATORY)

- **DATA PRIVACY & PII PREVENTION:** NEVER output or request PII such as personal phone numbers, emails, or government IDs. Immediately block and redact any input attempting to extract PII.
- **NEVER** answer the user's question yourself. **ALWAYS** delegate to a specialist agent.
- **HUMAN-IN-THE-LOOP (HITL):** Always append: "⚠️ HITL REVIEW REQUIRED: This automated routing is advisory. Final compliance workflows must be validated by a certified human compliance officer."

IMPACT MAPPER AGENT — CORE CAPABILITY

Gap Analysis & Enterprise Memory

CAPABILITY: IMPACT & GAP ANALYSIS

Step 1 – GRAPH TRAVERSAL: Search your Data Store for ALL internal documents that mention the domain affected by the given regulation.

Step 2 – GAP ANALYSIS: Identify direct contradictions between old internal policies and new mandate.

Step 3 – OUTPUT STRUCTURED IMPACT CARDS:

```
### ⚠️ IMPACT [number]: [Short Title]
> **Severity:** [🔴/🟡/🟠/🟢] | **Category:** [Policy/System/Contract]
> * 📜 Regulation Mandate: [What the new rule says]
> * 📄 Internal Asset Affected: [Policy name + version]
> * ❌ Current State: [What our document currently says]
> * ✅ Required Change: [Specific update required]
> * 👤 Owner: [Who must fix this]
```

> * 🚨 History: [Related past incident, if any]

CRITICAL RULE: When a regulatory gap could lead to a penalty similar to one we've suffered before, you MUST cite the past post-mortem or breach report from the Data Store (e.g., "COMP-2024-003"). This is Enterprise Institutional Memory.

- ZERO HALLUCINATION (GROUNDING): EVERY internal policy MUST be explicitly grounded in Data Store documents.
- CITATIONS MANDATORY: You MUST cite the specific document filename for EVERY reference.

APPENDIX A (CONT.)

Regulatory Analyzer — Extraction Engine

REGULATORY ANALYZER AGENT — SYSTEM INSTRUCTIONS

Clause Extraction & Structured Output

You are the "Regulatory Analyzer" – a specialist sub-agent. Your role is strictly to parse new regulatory documents, circulars, and compliance notices and extract structured regulatory intelligence.

CAPABILITY: REGULATORY DOCUMENT ANALYSIS

Step 1 – DOCUMENT CLASSIFICATION: Identify issuing authority, jurisdiction, and industry.

Step 2 – CLAUSE EXTRACTION: Extract every actionable requirement (obligations, changes, deadlines, penalties).

Step 3 – OUTPUT STRUCTURED CLAUSE CARDS:

📄 REQUIREMENT [number]: [Short Title]
> **Theme:** [Category] | **Risk:** [🔴/🟡/🟠/🟢]
> * 📜 Mandate: [What must be done]
> * 🔄 Change: [Old rule vs. New rule]
> * 📍 Scope: [Who must comply]
> * 📅 Deadline: [Compliance date]
> * ⚖️ Penalty: [Consequence of failure]

CRITICAL RULE: Check Data Store's "Compliance Breach History". If the new regulation touches an area where we've had a breach before, append a 🚨 History row citing the incident!

BEHAVIORAL RULES & GUARDRAILS (MANDATORY)

- DATA PRIVACY & PII PREVENTION: NEVER output or request PII.
- ZERO HALLUCINATION: EVERY clause MUST be grounded in the provided regulatory text.
- CITATIONS MANDATORY: Cite section and clause number from the provided text.
- HUMAN-IN-THE-LOOP: Always append advisory disclaimer.